

Online Algorithm

Presented by **Ting-An Chen**

Advisor: De-Nian Yang, Ming-Syan Chen

May 29, 2020

Outline

- Recap. Streaming v.s. Online algorithm
- Offline v.s. Online algorithm
- Online algorithm performance – competitiveness analysis
 - Case 1. Ski Rental Problem
 - Case 2. Deterministic Paging Problem
 - Case 3. Random Marking Algorithms
- Types of adversary

Recap. Streaming algorithm

Characteristics?

- Sequence of data
- Limited memory
- A sketch of data Summary
- Return output after seeing all data
- All data info. is known
- OPT is known
- → v.s. OPT approximation

What we want to know...

e.g. data properties –
[Recap.] “*Frequency moment*”
Zero-order: #distinct_values
1st -order: total input size

Compared with ... Online algorithm

- Sequence of data
- Limited memory
- A sketch of data
- Return output at **each** time stamp
- **Never know** the nature of the coming data
- OPT is **unknown**
- ➔ v.s. **Offline** OPT (known) [How to compare?](#)

Online algorithm (A) v.s. Offline OPT (OPT)

- Competitiveness analysis

- $\frac{\text{Cost}(A)}{\text{Cost}(\text{OPT})} \leq \text{bound}$, then A is *competitive*.

- [Def.] α – *competitive* online algorithm.

- σ : an input sequence

- c : a cost function

- \rightarrow

- A is said to be α – *competitive* if $c_A(\sigma) \leq \alpha \cdot c_{\text{OPT}}(\sigma)$.

- α : competitive ratio.

Case 1. Ski-rental problem

- Ski everyday
- Rent or buy the skiing equipment (daily decision)
 - Rent one day, \$1.
 - Buy, \$C.
- Assumption: might hurt each day then cannot ski.
- Let **d** be the total number of days skiing.
- Algorithm: “Rent for C days, then buy on (C+1)-th day.”
 - [pf.] **2** – *competitive* online algorithm, i.e., $c_A(\sigma) \leq \mathbf{2} \cdot c_{OPT}(\sigma)$.

case	$c_A(\sigma)$	$c_{OPT}(\sigma)$
If $d \leq C$	d	d
If $d > C$	2C	C

Online algorithm (A) v.s. Offline OPT (OPT)

- Competitiveness analysis

- $\frac{\text{Cost}(A)}{\text{Cost}(\text{OPT})} \leq \text{bound}$, then A is *competitive*.

Approximation
ratio?

- [Def.] α – *competitive* online algorithm.
- σ : an input sequence
- c : a cost function
- \rightarrow
- A is said to be α – *competitive* if $c_A(\sigma) \leq \alpha \cdot c_{\text{OPT}}(\sigma)$.
- α : competitive ratio.

Online v.s. Offline (traditional) algorithm

	Online	Offline
Compare to Offline OPT	Competitive ratio	Approximation ratio
Cost(A) related to... 1. Inputs	a. Unknown but <u>expected</u> b. Random w/o known patterns → Online alg. Cost(A): → fluctuate	a. Known b. Not random –OR– Random w/. Distribution → Offline alg. Cost(A): → stable
Cost(A) related to... 2. Algorithm	At different states... Same strategy - Deterministic Different strategies – Random	Single strategy
Inputs to the Alg.	Hard –OR– easy → average case	Hard → worst case

Studying worst case in Online algorithm?

- **Adversary!!** To consider the case: the inputs make the algorithm worst
→ competitive ratio
- Known:
 - 1. the online algorithm, i.e., the strategy at each output state
 - 2. the outcomes of events
 - 3. inputs in the past
- Unknown:
 - 1. inputs in the future

Adversary

- Generate the input sequence that may induce the worst case performance

Oblivious The adversary designs the input sequence σ at the beginning. It does not know any randomness used by algorithm \mathcal{A} .

Adaptive At each time step t , the adversary knows all randomness used by algorithm \mathcal{A} thus far. In particular, it knows the exact state of the algorithm. With these in mind, it then picks the $(t + 1)$ -th element in the input sequence.

Fully adaptive The adversary knows all possible randomness that will be used by the algorithm \mathcal{A} when running on the full input sequence σ . For instance, assume the adversary has access to the same pseudorandom number generator used by \mathcal{A} and can invoke it arbitrarily many times while designing the adversarial input sequence σ .

Online v.s. Offline (traditional) algorithm

Adversarial
inputs?

Random
inputs

	Online	Offline
Compare to Offline OPT	Competitive ratio	Approximation ratio
Cost(A) related to... 1. Inputs	a. Unknown but <u>expected</u> b. Random w/o known patterns → Online alg. Cost(A): → fluctuate	a. Known b. Not random –OR– Random w/. Distribution → Offline alg. Cost(A): → stable
Cost(A) related to... 2. Algorithm	At different states... Same strategy - Deterministic Different strategies – Random	Single strategy
Inputs to the Alg.	Hard –OR– easy → average case	Hard → worst case

Online algorithm – random v.s. adversarial inputs

	Online (random inputs)	Online (adversarial inputs)	Offline
Compare to Offline OPT	Competitive ratio	Competitive ratio	Approximation ratio
Cost(A) related to... 1. Inputs	a. Unknown but <u>expected</u> b. Random w/o known patterns → Online alg. Cost(A): → fluctuate	a. Known (simulated) b. Not random → Online alg. Cost(A): → Stable → Online alg. $\text{Cost}(A) \leq \text{Cost}(A)'$	a. Known b. Not random –OR– Random w/. Distribution → Offline alg. Cost(A): → stable
Inputs to the Alg.	Hard –OR– easy → average case	Hard → worst case	Hard → worst case

Example. Sorting

- Offline – Selection sort

Initial

1	5	2	3
---	---	---	---

Find 1st min.

1	5	2	3
---	---	---	---

Find 2nd min.

1	2	5	3
---	---	---	---

...

Final

1	2	3	5
---	---	---	---

- Online – Insertion sort

Initial

1

Comes 1st one, compare and insert

5

1	5
---	---

Comes 2nd one, compare and insert

2

1	2	5
---	---	---

...

Final

1	2	3	5
---	---	---	---

Case 2. Paging problem

- Hard disk – large memory, slow access
- Cache – small memory, fast access
- A sequence of page requests (from cache)
- *Page fault* if requested info. is not in cache
- → access from hard disk
- → large access costs
- **Problem:**
 - what data is to be stored in cache s.t. fewest page faults
 - more precisely, **which data in cache is to be evicted when a new data is requested**

最不近(最久之前) request 的，先 evict 拔除

Paging problem – Least Recently Used (LRU)

- Example

request	cache elements	page fault	evicted item
a	-, -, -	True	-
b	a, -, -	True	-
c	a, b, -	True	-
d	a, b, c	True	a
a	d, b, c	True	b
e	d, a, c	True	c
b	d, a, e	True	d
a	b, a, e	False	
c	b, a, e	True	e
e	b, a, c	True	b

Deterministic

Online

Algorithm:

With specific
strategy

Paging problem

- Claim: If A is a deterministic online algorithm that is *α – competitive*, then $\alpha \geq k$, where k is the cache size. (at most k pages in cache), and total $(k+1)$ distinct pages.
 - [pf.]

$c_A(\sigma)$	$c_{OPT}(\sigma)$
Worst case $= \sigma $	

All requests are faults.

Evict p_i at t , then request p_i at $(t+1)$.

Paging problem

- [pf.]

A k-page sized cache




$\sigma(t)$

$K+1$

$\sigma(t+1)$

$\sigma(t+k-1)$



$c_A(\sigma)$	$c_{OPT}(\sigma)$
Worst case $= \sigma $	

- Suppose that OPT has a fault on some request $\sigma(t)$.
- OPT can evict a page is not requested during the next $k-1$ requests $\sigma(t+1), \dots, \sigma(t+k-1)$.
- Thus, on any k consecutive requests, OPT has ≤ 1 fault

Paging problem

- Claim: If A is a deterministic online algorithm that is *α – competitive*, then $\alpha \geq k$, where k is the cache size. (at most k pages in cache), and total $(k+1)$ distinct pages.
 - [pf.]

$c_A(\sigma)$	$c_{OPT}(\sigma)$
Worst case $= \sigma $	$\leq \frac{ \sigma }{k}$

Every k pages
has ≤ 1 fault

Deterministic v.s. Randomized

- Case 3. Random Marking Algorithm (RMA)
 - Initialize all pages as marked
 - Upon request of a page p
 - If p is not in cache,
 - * If all pages in cache are marked, unmark all
 - * Evict a random unmarked page
 - Mark page p

Deterministic v.s. Randomized

- Case 3. Random Marking Algorithm (RMA)

Example Suppose $k = 3$, $\sigma = (2, 5, 2, 1, 3)$.

Suppose the cache is initially:

Cache	1	3	4
Marked?	✓	✓	✓

When $\sigma(1) = 2$ arrives, all pages were unmarked. Suppose the random eviction chose page '3'. The newly added page '2' is then marked.

Cache	1	2	4
Marked?	✗	✓	✗

When $\sigma(2) = 5$ arrives, suppose random eviction chose page '4' (between pages '1' and '4'). The newly added page '5' is then marked.

Cache	1	2	5
Marked?	✗	✓	✓

When $\sigma(3) = 2$ arrives, page '2' in the cache is marked (no change).

Cache	1	2	5
Marked?	✗	✓	✓

When $\sigma(4) = 1$ arrives, page '1' in the cache is marked. At this point, any page request that is not from $\{1, 2, 5\}$ will cause a full unmarking of all pages in the cache.

Cache	1	2	5
Marked?	✓	✓	✓

Deterministic v.s. Randomized

- Competitiveness analysis
- A: c-competitive

$$E[C_A(\sigma)] \leq c \cdot C_{OPT}(\sigma)$$

Summary

- Recap. Streaming algorithm v.s. Online algorithm
- Online v.s. Offline
 - Example. sorting
- Online performance – competitiveness analysis
 - Case 1. Ski Rental Problem
 - Case 2. Deterministic Paging Problem
 - Case 3. Random Marking Algorithms
- Types of adversary

	Online (random inputs)	Online (adversarial inputs)	Offline
Compare to Offline OPT	Competitive ratio	\leq Competitive ratio	Approximation ratio
Cost(A) related to... 1. Inputs	a. Unknown but <u>expected</u> b. Random w/o known patterns → Online alg. Cost(A): → fluctuate	a. Known (simulated) b. Not random → Online alg. Cost(A)': → Stable → Online alg. $\text{Cost(A)} \leq \text{Cost(A)'}'$	a. Known b. Not random –OR– Random w/. Distribution → Offline alg. Cost(A): → stable
Cost(A) related to... 2. Algorithm	At different states... Same strategy - Deterministic Different strategies – Random	At different states... Same strategy - Deterministic Different strategies – Random	Single strategy
Inputs to the Alg.	Hard –OR- easy → average case	Hard → worst case	Hard → worst case