


# Online Algorithms with ORACLE and ML Advice (and why these are different)

Joint Study Group, 20200709

# Timeline

Traditional competitive analysis  
worst-case analysis



Spielman&Teng  
"worst-case inputs  
are brittle"

2004:  
smoothed analysis

"self-improving" algorithms:  
learn the input distribution

2011

Assuming inputs  
are stochastic  
(follow distribution;  
statistical properties)

Unrealistic; hard to verify

At least always better  
than naïve worst-case  
without advice

Cole&Roughgarden  
"alg cannot see every input,  
but can sample  
from them"

2014:  
learning from  
sample

2015  
Full-spectrum  
robustness

2018  
(This ICML)  
Provide foundation  
of this paradigm:  
treat ML as a black-box

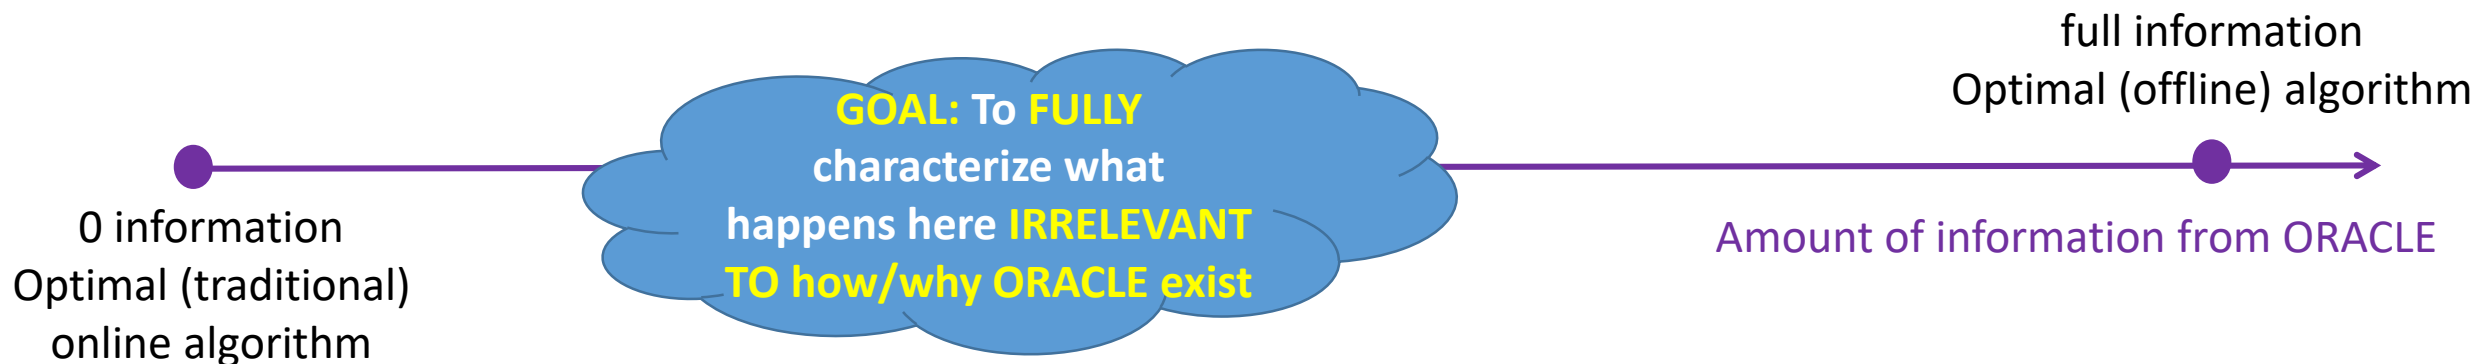
Assuming a perfect oracle  
Minimize the oracle queries  
needed for a performance  
requirement

# Online Alg with PERFECT ORACLE Advice

Joan Boyar, Lene M. Favrholdt, Christian Kudahl, Kim S. Larsen,  
Jesper W. Mikkelsen: Online Algorithms with Advice: A Survey.  
SIGACT News 47(3): 93-129 (2016)

- Input is adversarial
- Decisions are irrevocable
- Algorithm knows ~~nothing~~ about the inputs before arrival

some information bits from ORACLE



1. How many bits of advice are necessary and sufficient to obtain a competitive ratio  $c$ ?
2. Given a fixed  $B$  bits of advice, how good can  $c$  be?

# What exactly is “fully characterize”?




Figure 1: The (asymptotic) trade-off between competitive ratio and advice for PAGING. The function  $f(k)$  is a rapidly growing function of  $k$  (but does not depend on  $n$ ). Consider a trade-off point  $(b, c)$  where  $b$  is a number of advice bits and  $c$  is a competitive ratio. The red area shows those trade-offs which provably cannot be achieved. The green area shows those trade-offs that we currently have algorithms achieving. It is an open problem if trade-offs in the white area are achievable or not. The horizontal dashed lines are the best possible competitive ratios of deterministic and randomized algorithms without advice.

# Some insights for PERFECT ORACLE

- Connection between randomized online algorithms
  - Amount of advice  $\approx$  amount of randomness needed (“guessing” advice)
- De-online
  - Use an online algorithm with  $b$  bits of advice in time  $O(T(n))$
  - Obtain an offline approximation algorithm in time  $O(2^b T(n))$
- Leads to rigorous “Complexity Classes” for online problems
  - “advice complexity”

# Online optimization with imperfect ML advice

- Key difference: These ML-based “predictors” are NOT ORACLES because they are estimators
- Want to show:  
incorporating the advice enhances the quality of optimization task  
with the quality of enhancement depending on the quality of advice
  - Robustness
    - At least as good as naïve online algorithm, if predictor is bad
  - Consistency
    - At least as good as optimal offline algorithm, if predictor is good/powerful

# Ski Rental and Non-Clairvoyant Job Scheduling with ML advice

Manish Purohit, Zoya Svitkina, Ravi Kumar: Improving Online Algorithms via ML Predictions. NeurIPS 2018: 9684-9693

Sreenivas Gollapudi, Debmalya Panigrahi: Online Algorithms for Rent-Or-Buy with Expert Advice. ICML 2019: 2319-2327

- Traditional task: **Rent or Buy?**

- Deterministic Break-even algorithm: 2-competitive (optimal)
- Randomized algorithm:  $\frac{e}{e-1} \approx 1.58$ -competitive (optimal)

$x$ : actual # of skiing days

$y$ : predicted # of skiing days

$\eta = |y - x|$ : prediction error

- Traditional task: **Task Scheduling (jobs/machines)**

- With uncertainty: running time of job unknown until finishes (
- Deterministic Round-Robin algorithm: 2-competitive (optimal)

$x_i$ : actual running time of task  $i$

$y_i$ : predicted running time of task  $i$

$\eta_i = |y_i - x_i|$ : prediction error of task  $i$

Let  $\lambda \in (0, 1)$  be a hyperparameter. For the ski rental problem with a predictor, we first obtain a deterministic online algorithm that is  $(1 + 1/\lambda)$ -robust and  $(1 + \lambda)$ -consistent (Section 2.2). We next improve these bounds by obtaining a randomized algorithm that is  $(\frac{1}{1 - e^{-(\lambda - 1/b)}})$ -robust and  $(\frac{\lambda}{1 - e^{-\lambda}})$ -consistent, where  $b$  is the cost of buying (Section 2.3). For the non-clairvoyant scheduling problem, we obtain a randomized algorithm that is  $(2/(1 - \lambda))$ -robust and  $(1/\lambda)$ -consistent. Note that the consistency bounds for all these algorithms circumvent the lower bounds, which is possible only because of the predictions.

Q: how about multiple (independent) predictors?

# Online Weighted Paging + ML

**Theorem 1.1.** *For weighted paging with PRP, any deterministic algorithm is  $\Omega(k)$ -competitive, and any randomized algorithm is  $\Omega(\log k)$ -competitive.*

**Theorem 1.2.** *For weighted paging with  $\ell$ -strong lookahead where  $\ell \leq n - k$ , any deterministic algorithm is  $\Omega(k)$ -competitive, and any randomized algorithm is  $\Omega(\log k)$ -competitive.*

**SPRP (“strong per-request prediction”):** On a request for page  $p$ , the predictor gives the next time-step when  $p$  will be requested *and all page requests till that request.*

**Theorem 1.3.** *There is a deterministic 2-competitive for weighted paging with SPRP.*

**Theorem 1.4.** *For weighted paging with SPRP, there is no deterministic algorithm whose cost is  $o(k) \cdot \text{OPT} + o(\ell_{pd})$ , and there is no randomized algorithm whose cost is  $o(\log k) \cdot \text{OPT} + o(\ell_{pd})$ .*

- Propose a new suitable model for weighted paging
- Design new online algorithms for new model
- Show lower bounds for new model



# And More...

- Bin Packing/List Update with both ORACLE and untrusted ML advice
  - Spyros Angelopoulos, Christoph Dürr, Shendan Jin, Shahin Kamali, Marc P. Renault: Online Computation with Untrusted Advice. ITCS 2020: 52:1-52:15

# Sidenote: Optimization with ML advice

- Difference: the optimization task is **NOT of sequential/online flavor**
- Still, there is an **optimization-like problem**, for which one needs to **learn/predict some key information** to achieve good results
- *Focusing on the optimization part*, show good results *assuming a black-box predictor exists*
- Example: Andres Muñoz Medina, **Sergei Vassilvitskii**: **Revenue Optimization with Approximate Bid Predictions**. NIPS 2017: 1858-1866

# Revenue Optimization with Approximate Bid Predictions

- Traditional task: **Setting a good reserve price in auctions**
  - Traditionally: **easy for known-distributional-bidders [Myerson 81]**
  - Challenge 1: **distribution not known (need to sample)**
  - Challenge 2: **item is highly heterogeneous (no history/data to learn from)**
- Why not directly ML? -> **Revenue function is highly non-continuous/non-convex**

**Our results.** We show that given a predictor of the bid with squared loss of  $\eta^2$ , we can construct a reserve function  $r$  that extracts all but  $g(\eta)$  revenue, for a simple increasing function  $g$ . (See

$$\Phi(\mathcal{C}) := \sum_{j=1}^k \sqrt{\sum_{i,i': x_i, x_{i'} \in C_k} (b_i - b_{i'})^2} = 2 \sum_{j=1}^k m_j \hat{\sigma}_j \quad (2)$$

**Theorem 2.** Let  $\delta > 0$  and let  $r_k$  denote the output of Algorithm 1 then  $r_k \in G(h, k)$  and with probability at least  $1 - \delta$  over the samples  $\mathcal{S}$ :

$$\hat{S}(r_k) \leq (3\hat{B})^{1/3} \left( \frac{1}{2m} \Phi(\mathcal{C}^h) \right) \leq (3\hat{B})^{1/3} \left( \frac{1}{2k} + 2 \left( \eta^2 + \sqrt{\frac{\log 1/\delta}{2m}} \right)^{1/2} \right)^{2/3}.$$